

Hybrid classical-quantum learning

www.quantum.ml

PATRICK STECHER

Technical University Munich

patrick.stecher@tum.de

2020-06-18

Introduction


Architecture

Results


Conclusion & Outlook

Section 1

Introduction




¹[Hubregtsen et al, 2020]



Goal: Classify this Data set!¹

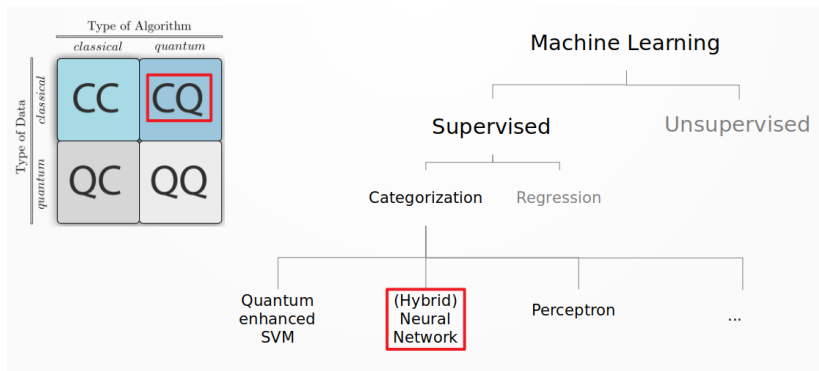
¹[Hubregtsen et al, 2020]



Goal: Classify this Data set!¹ (Using a quantum computer 😊)

¹[Hubregtsen et al, 2020]

30.000ft overview



Definition(s)

No fixed definition for "Hybrid":

- ▶ Using classical data on a quantum circuit (or the other way around)

Definition(s)

No fixed definition for "Hybrid":

- ▶ Using classical data on a quantum circuit (or the other way around)
- ▶ Down projection of high-dimensional inputs, (e.g. pictures) to low-dimensional representations (e.g. using convolutional NN), suitable for running on a quantum computer typically consisting of a single-digit number of qubits.

Definition(s)

No fixed definition for "Hybrid":

- ▶ Using classical data on a quantum circuit (or the other way around)
- ▶ Down projection of high-dimensional inputs, (e.g. pictures) to low-dimensional representations (e.g. using convolutional NN), suitable for running on a quantum computer typically consisting of a single-digit number of qubits.
- ▶ **Optimization of a Parameterized Quantum Circuit (PQC) using classical optimization techniques such as gradient descent.**

Section 2

Architecture

Architecture

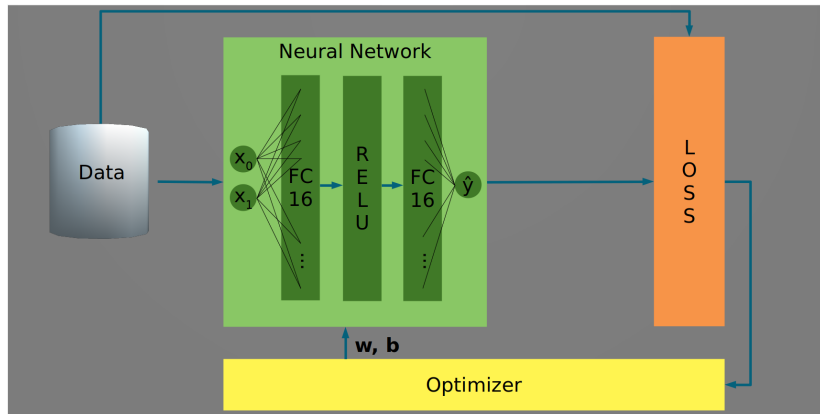


Figure: architecture of a classical network

Architecture

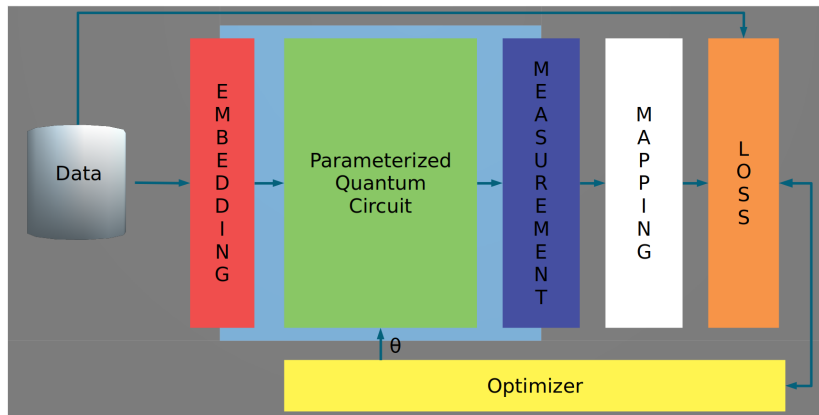


Figure: General architecture of a hybrid-classical network

The Data

- ▶ 9 synthetic data sets



The Data

- ▶ 9 synthetic data sets
- ▶ 1500 samples each




The Data

- ▶ 9 synthetic data sets
- ▶ 1500 samples each
- ▶ different "difficulties"



The Data

- ▶ 9 synthetic data sets
- ▶ 1500 samples each
- ▶ different "difficulties"



Embedding



How to get the data into the Quantum circuit?

- ▶ A mapping function $f : \mathbf{x} \rightarrow U_{\phi(\mathbf{x})}|0\rangle^{\otimes n}$ (for n qubits) is needed.
- ▶ a feature map² $\phi(\mathbf{x})$ is needed.

²e.g.: $\phi : \mathbf{x} \rightarrow [0, 2\pi)$

³[Benedetti et al, 2019]

Embedding



How to get the data into the Quantum circuit?

- ▶ A mapping function $f : \mathbf{x} \rightarrow U_{\phi(\mathbf{x})}|0\rangle^{\otimes n}$ (for n qubits) is needed.
- ▶ a feature map² $\phi(\mathbf{x})$ is needed.

Embedding scheme	Formula	Kernel ³
Amplitude	$x \in \{0, 1\}^n \rightarrow x\rangle$	linear kernel
Product	$x \in \mathbb{R}^N \rightarrow \psi_x\rangle = \cos(x_j) 0\rangle + \sin(x_j) 1\rangle$	highly non-linear kernels

²e.g.: $\phi : \mathbf{x} \rightarrow [0, 2\pi)$

³[Benedetti et al, 2019]

Embedding Circuit





Figure: Embedding circuit


Parameterized Quantum Circuit



- ▶ The basic idea of quantum computing is similar to that of kernel methods in machine learning, to efficiently perform computations in an intractably large Hilbert/Kernel space.

⁴[Sim et al. 2019]


Parameterized Quantum Circuit



- ▶ The basic idea of quantum computing is similar to that of kernel methods in machine learning, to efficiently perform computations in an intractably large Hilbert/Kernel space.
- ▶ Parameterized Quantum Circuits are a combination of multiple quantum gates operating on one or multiple qubits. The quantum gates have free parameters which can be optimized to fit a desired probability distribution.

⁴[Sim et al. 2019]


Parameterized Quantum Circuit





- ▶ The basic idea of quantum computing is similar to that of kernel methods in machine learning, to efficiently perform computations in an intractably large Hilbert/Kernel space.
- ▶ Parameterized Quantum Circuits are a combination of multiple quantum gates operating on one or multiple qubits. The quantum gates have free parameters which can be optimized to fit a desired probability distribution.
- ▶ The circuits presented here and investigated in *Hubregtsen et al.* were originally proposed by *Sim et al.*⁴

⁴[Sim et al. 2019]

Parameterized Quantum Circuit




- ▶ The basic idea of quantum computing is similar to that of kernel methods in machine learning, to efficiently perform computations in an intractably large Hilbert/Kernel space.
- ▶ Parameterized Quantum Circuits are a combination of multiple quantum gates operating on one or multiple qubits. The quantum gates have free parameters which can be optimized to fit a desired probability distribution.
- ▶ The circuits presented here and investigated in *Hubregtsen et al.* were originally proposed by *Sim et al.*⁴




⁴[Sim et al. 2019]

Measurement

- ▶ The four qubits are then measured in a fixed but arbitrary basis. Arbitrary since you can always rotate into a different basis, in this case this rotation is just learned.




Measurement




- ▶ The four qubits are then measured in a fixed but arbitrary basis. Arbitrary since you can always rotate into a different basis, in this case this rotation is just learned.
- ▶ The process of preparing quantum states, embedding data, applying rotations according to the parameter vector θ and measuring the outcome are performed multiple times.

Measurement




- ▶ The four qubits are then measured in a fixed but arbitrary basis. Arbitrary since you can always rotate into a different basis, in this case this rotation is just learned.
- ▶ The process of preparing quantum states, embedding data, applying rotations according to the parameter vector θ and measuring the outcome are performed multiple times.
- ▶ This is referred to as “shots” and will lead to a probability distribution over the possible states.

Measurement




- ▶ The four qubits are then measured in a fixed but arbitrary basis. Arbitrary since you can always rotate into a different basis, in this case this rotation is just learned.
- ▶ The process of preparing quantum states, embedding data, applying rotations according to the parameter vector θ and measuring the outcome are performed multiple times.
- ▶ This is referred to as “shots” and will lead to a probability distribution over the possible states.




Mapping function

- ▶ How to get class labels from a distribution?




Mapping function

- ▶ How to get class labels from a distribution?
- ▶ When measuring n qubits: $j \in \{0, 1\}^n$, some mapping function $j \rightarrow \{0, 1\}$ is needed.




Mapping function

- ▶ How to get class labels from a distribution?
- ▶ When measuring n qubits: $j \in \{0, 1\}^n$, some mapping function $j \rightarrow \{0, 1\}$ is needed.
- ▶ Big hurdle while implementing




Mapping function

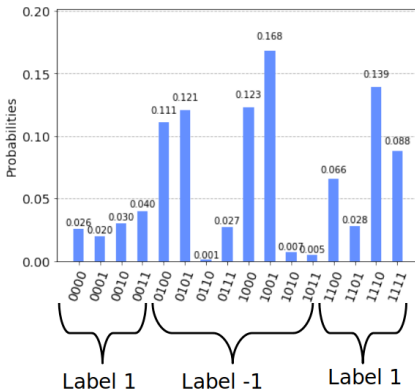
- ▶ How to get class labels from a distribution?
- ▶ When measuring n qubits: $j \in \{0, 1\}^n$, some mapping function $j \rightarrow \{0, 1\}$ is needed.
- ▶ Big hurdle while implementing
- ▶ $2^{2^4} = 65536$ possible binary mapping functions!




Mapping function



- ▶ How to get class labels from a distribution?
- ▶ When measuring n qubits: $j \in \{0, 1\}^n$, some mapping function $j \rightarrow \{0, 1\}$ is needed.
- ▶ Big hurdle while implementing
- ▶ $2^4 = 65536$ possible binary mapping functions!



Loss function



- ▶ The task of learning an arbitrary function from data is mathematically expressed as the minimization of a loss function $L(\theta)$, also known as the objective function, with respect to the parameter vector θ .


Loss function



- ▶ The task of learning an arbitrary function from data is mathematically expressed as the minimization of a loss function $L(\theta)$, also known as the objective function, with respect to the parameter vector θ .
- ▶ Squared Loss:

$$L(\theta) = (y - \hat{y})^2$$

Loss function




- ▶ The task of learning an arbitrary function from data is mathematically expressed as the minimization of a loss function $L(\theta)$, also known as the objective function, with respect to the parameter vector θ .
- ▶ Squared Loss:

$$L(\theta) = (y - \hat{y})^2$$




$$\hat{y} = label * max(probabilities)$$

Optimizer




- ▶ Goal: minimize loss function

Optimizer



- ▶ Goal: minimize loss function
- ▶ Problem: Backpropagation is no easy feat


Optimizer



- ▶ Goal: minimize loss function
- ▶ Problem: Backpropagation is no easy feat
- ▶ Solution: Parameter shift rule:


$$\frac{\partial L}{\partial \theta_j} \approx \frac{L(\boldsymbol{\theta} + c\boldsymbol{\Delta}) - L(\boldsymbol{\theta} - c\boldsymbol{\Delta})}{2c\Delta_j}$$

Optimizer



- ▶ Goal: minimize loss function
- ▶ Problem: Backpropagation is no easy feat
- ▶ Solution ???


Optimizer



- ▶ Goal: minimize loss function
- ▶ Problem: Backpropagation is no easy feat
- ▶ Solution ???
- ▶ Has to be executed twice for each parameter!


$$2 * \#Params * \#Shots$$

Optimizer




- Solution: SPSA (Simultaneous perturbation stochastic approximation)

Optimizer



- ▶ Solution: SPSA (Simultaneous perturbation stochastic approximation)
- ▶ only $2 * \#Shots$ evaluations


Optimizer



- ▶ Solution: SPSA (Simultaneous perturbation stochastic approximation)
- ▶ only $2 * \#Shots$ evaluations

$$\frac{\partial L}{\partial \theta_j} \approx \frac{L(\theta + \Delta \mathbf{e}_j) - L(\theta - \Delta \mathbf{e}_j)}{2\Delta}$$

Optimizer




- ▶ Solution: SPSA (Simultaneous perturbation stochastic approximation)
- ▶ only $2 * \#Shots$ evaluations

$$\frac{\partial L}{\partial \theta_j} \approx \frac{L(\theta + \Delta \mathbf{e}_j) - L(\theta - \Delta \mathbf{e}_j)}{2\Delta}$$


- ▶ → evaluate the gradient for each dimension joint, not independently.

Optimizer



- Now, that we have a gradient, gradient descent can be used!

Optimizer



- ▶ Now, that we have a gradient, gradient descent can be used!
- ▶ Or, even better ADAM, which has some improvements over regular gradient descent

Section 3

Results

- ▶ I used pennyLane by XANADU.

P E N N Y L A N E



- ▶ I used pennyLane by XANADU.

P E N N Y L A N E



- ▶ python framework

- ▶ I used pennyLane by XANADU.

P E N N Y L A N E  XANADU

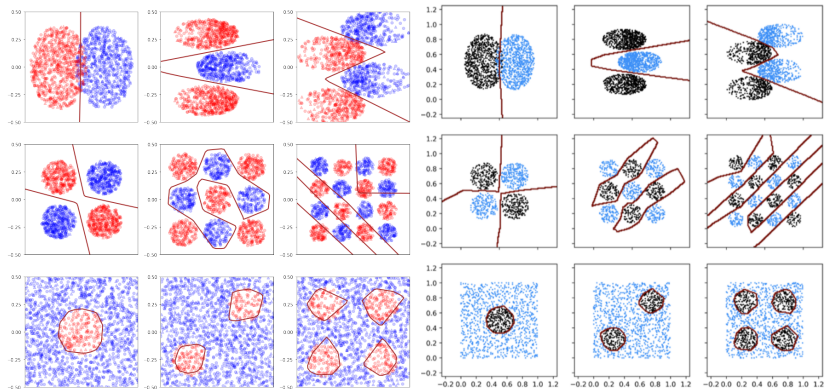
- ▶ python framework
- ▶ focused on quantum machine learning and quantum chemistry

- ▶ I used pennyLane by XANADU.

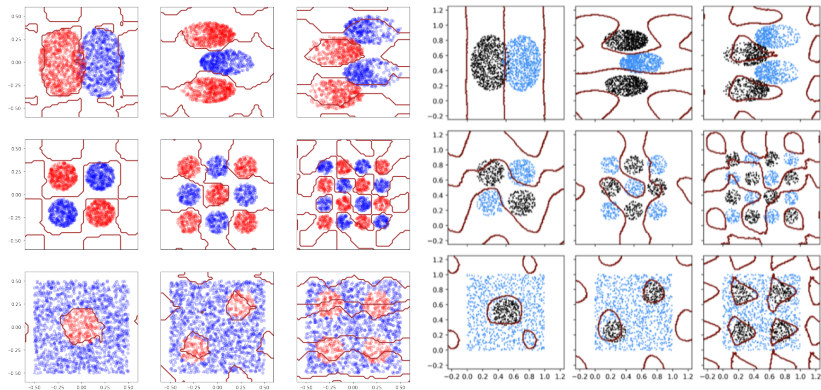


- ▶ python framework
- ▶ focused on quantum machine learning and quantum chemistry
- ▶ Interfaces with Pytorch, Tensorflow and even Qiskit.



Results I: Neural Networks



Results II: Quantum Circuits



Results III: QC vs. NN



Section 4

Conclusion & Outlook

Conclusion

- ▶ Hybrid learning will be a powerful technique on near term noisy quantum computers
- ▶ To date, all supervised learning experiments involved scaled-down, often trivial, data sets due to the limitation of available quantum hardware, and demonstrations at a more realistic scale are desirable.
- ▶ One possible pitfall is that as the circuits become more expressive, the optimization landscape might also become harder to explore.
- ▶ Provided that we can efficiently load or prepare quantum data in a qubit register, PQC models will deliver a clear advantage over classical methods for quantum learning tasks.

Outlook

What's next?

- ▶ Have a look at the code and try it yourself!
- ▶ Try a different initialization scheme to avoid barren plateaus
- ▶ Try combinations of different circuits, layers, embedding or cost functions!
- ▶ Try different optimizers, such as gradient free optimizer (e.g. genetic algorithms) or try an analytical gradient!
- ▶ Try regularization!
- ▶ Try to find a real-world data set!
- ▶ Go deeper!
- ▶ read the documentation in-depth at www.qu.antum.ml

Thank you for your attention!
:)

Training scheme

- ▶ Another problem; as only the Learning rate schedule parameters are mentioned:

Training scheme

- ▶ Another problem; as only the Learning rate schedule parameters are mentioned:
- ▶ Learning rate schedule: changed every 30 epochs (4, 3, 2, 1.5, 1, 0.5)

Training scheme

- ▶ Another problem; as only the Learning rate schedule parameters are mentioned:
- ▶ Learning rate schedule: changed every 30 epochs (4, 3, 2, 1.5, 1, 0.5)
- ▶ Minibatch size, betas, learning rate decay, ... are unknown (and were determined by trial& error)

Backup: Results IV: accuracy scores

Setup		Performance									
Type	Alg.	1a	1b	1c	2a	2b	2c	3a	3b	3c	Average
Classical	NN – 2l	96%	100%	100%	100%	99%	98%	100%	99%	97%	99%
Classical (mine)	NN – 2l	95%	100%	100%	100%	96%	87%	100%	98%	94%	97%
Hybrid	c19 – 2l	95%	97%	84%	91%	94%	74%	93%	91%	87%	90%
Hybrid (mine)	c19 – 2l	95%	99%	98%	100%	99%	98%	96%	97%	80%	95%

Sources



Hubregtsen, T. and Pichelmaier, J and Bertels, K. (2020).
"Evaluation of Parameterized Quantum Circuits: on the design, and the relation between classification accuracy, expressibility and entangling capability"
[arXiv preprint arXiv:2003.09887](#).



Benedetti, M. and Lloyd, E. and Sack, S. and Fiorentini, M. (2019).
"Parameterized quantum circuits as machine learning models"
[arXiv preprint arXiv:1906.07682v2](#).



Sim, S. and Johnson, PD. and Aspuru-Guzik, A. (2019).
"Expressibility and entangling capability of parameterized quantum circuits for hybrid quantum-classical algorithms"
[arXiv preprint arXiv:1906.07682v2](#).